

PROBLEM STATEMENT:-

TO PREDICT THE RAIN FALL BASED ON VARIOUS FEATURES OF THE DATASET

Importing All The Required Libraries

In [1]:

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
import matplotlib.pyplot as plt
import seaborn as sns
```

STEP-1:Data Collection

In [2]:

```
df=pd.read_csv(r"C:\Users\thara\Downloads\RainFall.csv")
df
```

Out[2]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	A
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	27
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	42
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	46
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	42
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	71
...
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	52
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	63
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	35
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	59
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	21

641 rows × 19 columns



STEP-2:Data Cleaning And Preprocessing

In [3]:

```
df.head()
```

Out[3]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0

In [4]:

```
df.tail()
```

Out[4]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AU
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217

In [5]:

df.describe()

Out[5]:

	JAN	FEB	MAR	APR	MAY	JUN	JUL
count	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000
mean	18.355070	20.984399	30.034789	45.543214	81.535101	196.007332	326.033697
std	21.082806	27.729596	45.451082	71.556279	111.960390	196.556284	221.364643
min	0.000000	0.000000	0.000000	0.000000	0.900000	3.800000	11.600000
25%	6.900000	7.000000	7.000000	5.000000	12.100000	68.800000	206.400000
50%	13.300000	12.300000	12.700000	15.100000	33.900000	131.900000	293.700000
75%	19.200000	24.100000	33.200000	48.300000	91.900000	226.600000	374.800000
max	144.500000	229.600000	367.900000	554.400000	733.700000	1476.200000	1820.900000

In [6]:

df.shape

Out[6]:

(641, 19)

In [7]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 641 entries, 0 to 640
Data columns (total 19 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   STATE_UT_NAME   641 non-null    object
 1   DISTRICT        641 non-null    object
 2   JAN             641 non-null    float64
 3   FEB             641 non-null    float64
 4   MAR             641 non-null    float64
 5   APR             641 non-null    float64
 6   MAY             641 non-null    float64
 7   JUN             641 non-null    float64
 8   JUL             641 non-null    float64
 9   AUG             641 non-null    float64
10   SEP             641 non-null    float64
11   OCT             641 non-null    float64
12   NOV             641 non-null    float64
13   DEC             641 non-null    float64
14   ANNUAL          641 non-null    float64
15   Jan-Feb         641 non-null    float64
16   Mar-May         641 non-null    float64
17   Jun-Sep         641 non-null    float64
18   Oct-Dec         641 non-null    float64
dtypes: float64(17), object(2)
memory usage: 95.3+ KB
```

In [8]:

```
df.isnull().any()
```

Out[8]:

```
STATE_UT_NAME    False
DISTRICT          False
JAN               False
FEB               False
MAR               False
APR               False
MAY               False
JUN               False
JUL               False
AUG               False
SEP               False
OCT               False
NOV               False
DEC               False
ANNUAL            False
Jan-Feb           False
Mar-May           False
Jun-Sep           False
Oct-Dec           False
dtype: bool
```

In [10]:

```
df.isnull().sum()
```

Out[10]:

```
STATE_UT_NAME    0
DISTRICT          0
JAN               0
FEB               0
MAR               0
APR               0
MAY               0
JUN               0
JUL               0
AUG               0
SEP               0
OCT               0
NOV               0
DEC               0
ANNUAL            0
Jan-Feb           0
Mar-May           0
Jun-Sep           0
Oct-Dec           0
dtype: int64
```

In [11]:

```
df.columns
```

Out[11]:

```
Index(['STATE_UT_NAME', 'DISTRICT', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN',  
      'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb',  
      'Mar-May', 'Jun-Sep', 'Oct-Dec'],  
      dtype='object')
```

In [12]:

```
df.duplicated().sum()
```

Out[12]:

```
0
```

In [13]:

```
df['ANNUAL'].value_counts()
```

Out[13]:

```
ANNUAL  
747.1      9  
2080.0     4  
1336.5     3  
1824.8     3  
2814.4     3  
..  
1037.6     1  
907.2      1  
944.5      1  
1003.3     1  
3253.1     1  
Name: count, Length: 591, dtype: int64
```

Feature Scaling:

To Split the data into training data and test data

In [14]:

```
x=df[["APR"]]  
y=df[["JAN"]]
```

In [15]:

```
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=50)
```


STEP-3: Data Visualization

In [16]:

```
slice = [12, 25, 50, 36, 19]
activities = ['DISTRICT', 'JAN', 'FEB', 'MAR', 'APR']
cols = ['r', 'b', 'c', 'g', 'orange']
plt.pie(slice, labels = activities, colors=cols, startangle=150, shadow=True, explode=(0, 0.1)
plt.title('Months')
```

Out[16]:

Text(0.5, 1.0, 'Months')




In [17]:

```
sns.pairplot(df)
```


Out[17]:

<seaborn.axisgrid.PairGrid at 0x2393a533bb0>



In [18]:

```
df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



STEP-4:- Data Modelling

Applying LINEAR REGRESSION

In [19]:

```
from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

13.76772905796956

Out[19]:

	coefficient
APR	0.08298

In [20]:

```
score=reg.score(X_test,y_test)
print(score)
```

0.13760810725167383

In [21]:


```
predictions=reg.predict(X_test)
```

In [22]:

```
plt.scatter(y_test,predictions)
```


Out[22]:

<matplotlib.collections.PathCollection at 0x2394ebfa8c0>



In [23]:

```
df500=df[:][:500]
sns.lmplot(x="APR",y="JAN",order=2,ci=None,data=df500)
plt.show()
```



In [24]:

```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
reg.fit(X_train,y_train)
reg.fit(X_test,y_test)
```

Out[24]:


LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [25]:

```
y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='ORANGE')
plt.plot(X_test,y_pred,color='BLUE')
plt.show()
```



In [26]:

```
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.11486832305419492

RIDGE MODEL

In [27]:

```
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

In [28]:

```
features= df.columns[0:5]
target= df.columns[-5]
```

In [29]:

```
x=np.array(df['APR']).reshape(-1,1)
y=np.array(df['JAN']).reshape(-1,1)
```

In [30]:

```
x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
```

In [31]:

```
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```

In [32]:

```
print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:


```
the train score for ridge model is0.9999999792491524
the test score for ridge model is0.9999999887465535
```

In [33]:

```
lr=LinearRegression()
```

In [34]:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red')
plt.plot(features,alpha=0.4,linestyle='none',marker="o",markersize=7,color='BLUE',label='LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



LASSO MODEL

In [35]:

```
print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.99912857000705


The test score for ls model is0.9991969731663574

In [36]:

```
pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="hist")
```

Out[36]:

<Axes: ylabel='Frequency'>



In [37]:

```

from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))

```

0.99999999999999198


0.9999999999999254

In [38]:

```

plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label='Lasso')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color="Aqua",label='LinearRegression')
plt.xticks(rotation=90)
plt.xticks(rotation = 90)
plt.legend()
plt.show()

```



ELASTICNET

In [39]:

```
from sklearn.linear_model import ElasticNet
eln=ElasticNet()
eln.fit(x,y)
print(eln.coef_)
print(eln.intercept_)
print(eln.score(x,y))
```

```
[ 9.93078025e-01  4.17330119e-03  0.00000000e+00 -2.56844715e-04
  4.33064019e-04]
0.04331617537314969
0.9999905490942288
```

In [45]:

```
y_pred_elastic = eln.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_error)
```

```
0.00427888506659882
```

Conclusion:

THE SCORE OF LINEAR REGRESSION IS :- 0.13760810725167383 THE SCORE OF RIDGE MODEL IS :- 0.9999999792491524 THE SCORE OF LASSO MODEL IS :- 0.9999999999999198 THE SCORE OF ELASTIC NET IS :- 0.9999905490942288 AMONG ALL MODELS LASSO YEILD HIGHEST ACCURACY.SO,WE PREFER LASSO MODEL FOR THIS DATA SET